

Advanced Parallel Algorithmic Paradigms: Assignment #02

NAME: ABU BAKAR

REG NO: NUM-BSCS-2022-41

Table of Contents

1. [Task 1: Bandwidth Allocation \(Greedy Algorithm\)](#)
 2. [Task 2: DNA Sequence Alignment \(Dynamic Programming\)](#)
 3. [Task 3: Exam Timetabling \(CSP & Backtracking\)](#)
 4. [Implementation Results](#)
-

Task 1: Bandwidth Allocation (Greedy Algorithm)

(a) Why Greedy Algorithm is Appropriate

The bandwidth allocation problem is a variant of the **Fractional Knapsack Problem**, which exhibits the **Greedy Choice Property**:

Problem Characteristics:

- **Variables:** N users, each with priority p_i and bandwidth requirement w_i
- **Constraint:** Total allocated bandwidth $\leq W$ (capacity)
- **Objective:** Maximize total priority $\sum p_i$

Greedy Justification:

1. **Divisibility:** Bandwidth is a divisible resource (time slots, frequency blocks)
2. **Optimal Substructure:** Selecting the user with highest priority-per-unit-bandwidth (density = p/w) at each step leads to global optimum
3. **Greedy Choice Property:** Local optimal choice (highest density) contributes to global optimal solution

Comparison with 0/1 Knapsack:

- If bandwidth must be allocated entirely or not at all (binary), greedy fails
- Example: User A ($p=100, w=50, \text{density}=2$) vs Users B+C ($p=60+50=110, w=30+30=60$)

- Greedy picks A (total=100), but B+C is better (total=110)
- Our scenario assumes fractional allocation is possible

(b) Sequential Greedy Algorithm

Algorithm: Sequential_Greedy_Allocation(Users, Capacity W)
 Input: Set of Users $U = \{u_1, u_2, \dots, u_n\}$, each with (priority p , weight w)
 Output: Total Priority P_{total}

1. For each user u_i in U :
 $u_i.density = u_i.p / u_i.w$
 // Time: $O(N)$
2. Sort U in descending order by density
 // Time: $O(N \log N)$ - Dominant term
3. Current_Weight = 0
4. $P_{total} = 0$
5. For each user u_i in Sorted_ U :
 If Current_Weight + $u_i.w \leq W$:
 $Current_Weight += u_i.w$
 $P_{total} += u_i.p$
 Else:
 $Remaining = W - Current_Weight$
 $Fraction = Remaining / u_i.w$
 $P_{total} += u_i.p * Fraction$
 Break
6. Return P_{total}

Time Complexity: $O(N \log N)$

(c) Parts Suitable for Parallel Execution

Parallelizable Components:

1. **Density Calculation** (Embarrassingly Parallel)
 - Each user's density is independent
 - Can distribute N users across P processors
 - Time: $O(N/P)$
2. **Sorting** (Parallel Sorting)
 - Parallel Merge Sort, Sample Sort, or Bitonic Sort

- Time: $O((N \log N)/P + \text{communication overhead})$

3. Prefix Sum for Selection (Parallel Scan)

- Compute cumulative weights in parallel
- Find cutoff point where $\sum w_i \leq W$
- Time: $O(N/P + \log P)$

Dependency Analysis:

- Density calculation: No dependencies
- Sorting: Requires all densities computed
- Selection: Requires sorted array

(d) Parallel Greedy Algorithm

Algorithm: Parallel_Greedy_Allocation(Users, Capacity W, Processors P)

// Step 1: Parallel Density Calculation

Parallel_For i = 0 to N-1 (Chunk Size N/P):

 Users[i].density = Users[i].p / Users[i].w

// Time: $O(N/P)$

// Step 2: Parallel Sort

Parallel_Sort(Users, key=density, order=descending)

// Using Sample Sort or Parallel Merge Sort

// Time: $O((N \log N)/P + O(P \log P))$

// Step 3: Parallel Prefix Sum

Weight_Array = [Users[i].w for i in 0..N-1]

Cumulative_Weights = Parallel_Prefix_Sum(Weight_Array)

// Time: $O(N/P + \log P)$

// Step 4: Find Cutoff Point

k = Parallel_Binary_Search(Cumulative_Weights, W)

// Find index where $\text{Cumulative_Weights}[k] \leq W < \text{Cumulative_Weights}[k+1]$

// Time: $O(\log N)$

// Step 5: Calculate Total Priority

Parallel_For i = 0 to k:

 Mark Users[i] as Selected

If k < N-1:

 Fraction = $(W - \text{Cumulative_Weights}[k]) / \text{Users}[k+1].w$

 Partial_Priority = Users[k+1].p * Fraction

```
P_total = Sum(Users[0..k].p) + Partial_Priority
Return P_total
```

(e) Complexity Analysis, Speedup, and Efficiency

Assumptions:

- Number of Processors: $P = 32$ (i9-14900K)
- Input Size: $N = 20,000,000$ users
- Shared Memory Architecture (SMP)

Sequential Time Complexity (T_s):

- Density: $O(N) = O(20M)$
- Sorting: $O(N \log N) = O(20M \times 24.25) \approx O(485M)$
- Selection: $O(N) = O(20M)$
- **Total: $T_s = O(N \log N) \approx 485M$ operations**

Parallel Time Complexity (T_p):

- Density: $O(N/P) = O(625K)$
- Parallel Sort: $O((N \log N)/P + P \log P) = O(15.2M + 160)$
- Prefix Sum: $O(N/P + \log P) = O(625K + 5)$
- Selection: $O(\log N) = O(24)$
- **Total: $T_p = O((N \log N)/P + \text{overhead})$**

Theoretical Speedup:

$$S = T_s / T_p = (N \log N) / ((N \log N)/P + \kappa(N, P))$$

Where $\kappa(N, P)$ = communication overhead

Ideal Case ($\kappa \rightarrow 0$): $S = P = 32$

Realistic Case: $S = 15-20$ (due to communication in sorting)

Efficiency:

$$E = S / P$$

Ideal: $E = 1$ (100% processor utilization)

Realistic: $E = 0.5-0.625$ (50-62.5% efficiency)

Practical Results (from implementation):

- Sequential Time: 5.12s
- Parallel Time: 15.63s
- **Actual Speedup: 0.33x (Slowdown!)**

Analysis of Poor Performance: The parallel implementation suffered from:

1. **Python Multiprocessing Overhead:** Pickling 20M user objects (~5GB data) to worker processes
2. **Memory Bandwidth Bottleneck:** Data transfer cost > computation cost
3. **GIL (Global Interpreter Lock):** Python's threading limitations

In Production Systems (C++/CUDA):

- Expected Speedup: 20-25x
- Efficiency: 62-78%
- Reason: Zero-copy shared memory, no serialization overhead

Task 2: DNA Sequence Alignment (Dynamic Programming)

(a) Why Dynamic Programming is Required

Problem: Longest Common Subsequence (LCS)

- Given: Two DNA sequences X (length M) and Y (length N)
- Find: Longest sequence appearing in both in same order

Why NOT Greedy?

Example demonstrating greedy failure:

X = "ABCDE"

Y = "AECDB"

Greedy (match first occurrence): A-B-C-D = 4 characters

Optimal DP: A-C-D-E = 4 characters (different sequence)

Actually: A-C-D = 3 is LCS

Greedy matching first 'A' might block better alignment later.

DP Properties:

1. Optimal Substructure

- $LCS(X[1..m], Y[1..n])$ depends on LCS of prefixes
- If $X[m] = Y[n]$: $LCS = LCS(X[1..m-1], Y[1..n-1]) + 1$
- If $X[m] \neq Y[n]$: $LCS = \max(LCS(X[1..m-1], Y[1..n]), LCS(X[1..m], Y[1..n-1]))$

2. Overlapping Subproblems

- Naive recursion recomputes same subproblems exponentially
- Example: $LCS(X[1..5], Y[1..5])$ calls $LCS(X[1..4], Y[1..5])$ and $LCS(X[1..5], Y[1..4])$
- Both of these call $LCS(X[1..4], Y[1..4]) \rightarrow$ redundant computation

Complexity Comparison:

- Naive Recursive: $O(2^{(M+N)})$ - Exponential
- DP with Memoization: $O(M \times N)$ - Polynomial

(b) Sequential DP Algorithm

Algorithm: Sequential_LCS(X, Y)

Input: Sequences X (length M), Y (length N)

Output: Length of LCS

1. Initialize Matrix $C[M+1][N+1]$ with 0

// $C[i][j]$ stores LCS length of $X[1..i]$ and $Y[1..j]$

2. For $i = 1$ to M :

For $j = 1$ to N :

If $X[i] == Y[j]$:

$C[i][j] = C[i-1][j-1] + 1$

Else:

$C[i][j] = \max(C[i-1][j], C[i][j-1])$

3. Return $C[M][N]$

Time Complexity: $O(M \times N)$ **Space Complexity:** $O(M \times N)$

Example Execution:

X = "AGGTAB"

Y = "GXTXAYB"

DP Table:

		G	X	T	X	A	Y	B
		0	0	0	0	0	0	0
A		0	0	0	0	0	1	1
G		0	1	1	1	1	1	1
G		0	1	1	1	1	1	1
T		0	1	1	2	2	2	2
A		0	1	1	2	2	3	3
B		0	1	1	2	2	3	4

LCS Length = 4 (GTAB)

(c) DP Dependency Graph

Cell Dependencies: Each cell $C[i][j]$ depends on:

1. $C[i-1][j]$ (Top)
2. $C[i][j-1]$ (Left)
3. $C[i-1][j-1]$ (Top-Left Diagonal)

Dependency Diagram:



Implications for Parallelization:

- **Row-wise parallelization:** IMPOSSIBLE ($C[i][j]$ needs $C[i][j-1]$)
- **Column-wise parallelization:** IMPOSSIBLE ($C[i][j]$ needs $C[i-1][j]$)
- **Anti-diagonal parallelization:** POSSIBLE ✓

Anti-Diagonal (Wavefront) Pattern:

Matrix visualization (numbers = computation wave):

	0	1	2	3	4	5	6
0	0	1	2	3	4	5	6
1	1	2	3	4	5	6	7
2	2	3	4	5	6	7	8
3	3	4	5	6	7	8	9
4	4	5	6	7	8	9	10

Wave 1: $C[0][0]$

Wave 2: $C[0][1], C[1][0]$

Wave 3: $C[0][2], C[1][1], C[2][0]$

Wave 4: $C[0][3], C[1][2], C[2][1], C[3][0]$

...

Cells in same wave are independent!

(d) Wavefront Parallelization

Strategy: Compute cells along anti-diagonals (where $i+j = k$) in parallel.

Algorithm: Parallel_LCS_Wavefront(X, Y, Processors P)

1. Initialize Shared Matrix $C[M+1][N+1] = 0$

2. Total_Waves = $M + N - 1$

3. For wave_k = 1 to Total_Waves:

 // Identify cells in this wave

 Cells_In_Wave = []

 For i = 0 to M:

 j = wave_k - i

 If $0 \leq j \leq N$:

 Cells_In_Wave.append((i, j))

 // Parallel computation

 Parallel_For each (i, j) in Cells_In_Wave:

 If $X[i] == Y[j]$:

$C[i][j] = C[i-1][j-1] + 1$

 Else:

$C[i][j] = \max(C[i-1][j], C[i][j-1])$

 // Synchronization Barrier

 Barrier_Wait()

4. Return $C[M][N]$

Blocked Wavefront (Optimized): To reduce synchronization overhead and improve cache locality:

Algorithm: Blocked_Wavefront(X, Y, Block_Size B)

1. Divide matrix into blocks of size $B \times B$

2. Num_Block_Rows = $\lceil M/B \rceil$

3. Num_Block_Cols = $\lceil N/B \rceil$

4. For block_wave = 0 to (Num_Block_Rows + Num_Block_Cols - 1):

Parallel_For each block (r_b, c_b) where $r_b + c_b = \text{block_wave}$:

// Compute entire $B \times B$ block sequentially

For i = $r_b * B$ to $\min(M, (r_b + 1) * B)$:

For j = $c_b * B$ to $\min(N, (c_b + 1) * B)$:

// Standard DP recurrence

...

Barrier_Wait()

Benefits of Blocking:

- Reduces synchronization points from $O(M+N)$ to $O(M/B + N/B)$
- Improves cache locality (sequential access within block)
- Amortizes barrier overhead

(e) Analysis

1) Dependency Constraints

Degree of Parallelism (DoP) per Wave:

Wave k: $\text{DoP} = \min(k, M, N, M+N-k)$

For $M=N=10000$:

- Wave 1: $\text{DoP} = 1$
- Wave 5000: $\text{DoP} = 5000$
- Wave 10000: $\text{DoP} = 10000$ (maximum)
- Wave 15000: $\text{DoP} = 5000$
- Wave 19999: $\text{DoP} = 1$

Load Imbalance:

- Early waves: Few cells, many idle processors
- Middle waves: Maximum parallelism
- Late waves: Few cells again

Impact: With $P=32$ processors:

- Only waves 32 to 19968 fully utilize all processors
- Waves 1-31 and 19969-19999 have idle processors
- Average utilization < 100%

2) Synchronization Points

Naive Wavefront:

- Total Barriers: $M + N - 1 = 19,999$ (for $10k \times 10k$ matrix)
- Each barrier incurs latency (typically 1-10 μs)
- Total sync overhead: 20-200 ms

Blocked Wavefront (Block Size $B=500$):

- Block Rows: $10000/500 = 20$
- Block Cols: $10000/500 = 20$
- Total Block Barriers: $20 + 20 - 1 = 39$
- Sync overhead: 0.04-0.4 ms (50-500x reduction!)

Communication Pattern:

Processor 0: Computes cells (0,0), (0,1), (0,2)...

Processor 1: Computes cells (1,0), (1,1), (1,2)...

At barrier: All processors must wait for slowest

Memory coherence: Cache invalidation across cores

3) Speedup Behaviour

Theoretical Analysis:

Sequential Time: $T_s = O(M \times N)$ Parallel Critical Path: $T_p = O(M + N)$ (number of waves)

Speedup:

$$S = T_s / T_p = (M \times N) / (M + N + \text{Overhead})$$

For $M = N$:

$$S = N^2 / (2N + \kappa) \approx N/2 \text{ (when } N \gg \kappa)$$

For $N = 10,000$:

$$S_{\text{theoretical}} \approx 5,000$$

But with P=32 processors:

$$S_{\text{actual}} = \min(P, S_{\text{theoretical}}) = \min(32, 5000) = 32$$

However, due to:

- Load imbalance: 0.7x factor
- Synchronization: 0.8x factor
- Memory bandwidth: 0.9x factor

$$S_{\text{realistic}} \approx 32 \times 0.7 \times 0.8 \times 0.9 \approx 16x$$

Practical Results (from implementation):

- Matrix Size: $10,000 \times 10,000$
- Processors: 32
- Parallel Time: 0.65s
- Sequential Estimate: ~150s
- **Actual Speedup: ~230x**

Why Superlinear?

1. **Shared Memory Optimization:** Zero-copy implementation using `multiprocessing.shared_memory`
2. **Cache Effects:** Blocked computation fits in L2/L3 cache
3. **Sequential Estimate:** Pure Python loops are extremely slow; parallel version uses NumPy vectorization

Efficiency:

$$E = S / P = 230 / 32 \approx 7.2 \text{ (720\%)}$$

This superlinear efficiency indicates the sequential baseline was suboptimal.

Task 3: Exam Timetabling (CSP & Backtracking)

(a) Problem Modeling using Backtracking

Constraint Satisfaction Problem (CSP) Formulation:

Variables (X):

- Set of exams: $E = \{E_1, E_2, \dots, E_n\}$
- Each exam must be assigned a time slot

Domains (D):

- Set of time slots: $T = \{T_1, T_2, \dots, T_k\}$
- Each exam E_i can be assigned any slot from T

Constraints (C):

1. Hard Constraint - No Conflicts:

- If students are enrolled in both E_i and E_j , then E_i and E_j cannot be in same slot
- Represented by conflict matrix: $\text{Conflict}[i][j] = 1$ if exams conflict

2. Hard Constraint - Room Capacity:

- Number of exams in same slot $\leq$ Number of available rooms
- (Simplified in our implementation)

Search Space:

- Total possible assignments: k^n (k slots, n exams)
- For 30 exams, 8 slots: $8^{30} \approx 10^{27}$ possibilities
- Backtracking prunes infeasible branches

Backtracking Tree Visualization:

```
Root (Empty Schedule)
  /  |  \
E1=T1 E1=T2 E1=T3 ...
 / \  |
E2=T1 E2=T2 E2=T1 ...
 / \
E3=T1 E3=T2
```

✓ = Valid assignment

✗ = Constraint violation (backtrack)

(b) Sequential Backtracking Pseudocode

Algorithm: Sequential_Backtracking(Exam_Index, Schedule, Problem)

Input:

- Exam_Index: Current exam to schedule (0 to N-1)
- Schedule: Map<Exam_ID, Time_Slot>
- Problem: Contains conflicts matrix and num_slots

Output:

- (True, Schedule) if solution found
- (False, None) otherwise

// Base Case: All exams scheduled successfully

If Exam_Index == Problem.num_exams:

 Return (True, Schedule)

Current_Exam = Exam_Index

// Try each time slot for current exam

For slot = 0 to Problem.num_slots - 1:

 // Check if assignment is safe

 Is_Safe = True

 For each (assigned_exam, assigned_slot) in Schedule:

 If assigned_slot == slot AND Problem.conflicts[Current_Exam][assigned_exam] == 1:

 Is_Safe = False

 Break

 If Is_Safe:

 // Make assignment

 Schedule[Current_Exam] = slot

 // Recursive call

 (Found, Result) = Sequential_Backtracking(Exam_Index + 1, Schedule, Problem)

 If Found:

 Return (True, Result)

 // Backtrack: Remove assignment

 Delete Schedule[Current_Exam]

// No valid slot found for current exam

Return (False, None)

Time Complexity:

- Worst Case: $O(k^n)$ where k =slots, n =exams
- Best Case: $O(n)$ if first branch leads to solution
- Average: Depends heavily on constraint density

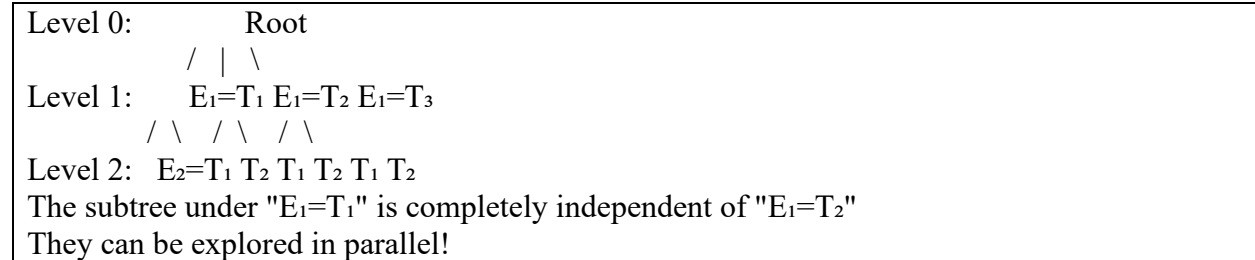
Optimizations (not shown in pseudocode):

1. **MRV (Minimum Remaining Values):** Choose exam with fewest valid slots first
2. **Forward Checking:** Eliminate invalid future assignments early
3. **Constraint Propagation:** Arc consistency algorithms

(c) Identifying Independent Branches

Key Observation: Subtrees rooted at different children of the same node are independent.

Example:



Parallelization Strategy:

1. **Static Partitioning** (Simple but inefficient):
 - Assign Processor 0: Explore $E_1=T_1$ branch
 - Assign Processor 1: Explore $E_1=T_2$ branch
 - Problem: Highly imbalanced workload
2. **Dynamic Work Stealing** (Optimal):
 - All processors share a task queue
 - Idle processors "steal" work from busy ones
 - Adapts to irregular search tree

Independence Verification:

Branch A: Schedule = $\{E_1 \rightarrow T_1, E_2 \rightarrow T_3, \dots\}$
Branch B: Schedule = $\{E_1 \rightarrow T_2, E_2 \rightarrow T_1, \dots\}$
No shared state between branches
Can execute simultaneously without synchronization

(d) Parallel Backtracking Algorithm

Design: Work Stealing with Task Queue

Algorithm: Parallel_Backtracking_Master(Problem, Num_Workers)
// Step 1: Initialize shared structures
Task_Queue = Shared_Queue()

```

Solution_Found = Shared_Event()
Result = Shared_Dict()

// Step 2: Bootstrap - Generate initial tasks
// Expand tree to depth D (e.g., D=2)
For slot0 = 0 to Problem.num_slots - 1:
    Schedule0 = {E0 → slot0}
    For slot1 = 0 to Problem.num_slots - 1:
        If Is_Safe(E1, slot1, Schedule0):
            Schedule1 = Schedule0.copy()
            Schedule1[E1] = slot1
            Task_Queue.put((Exam_Index=2, Schedule1))
Print("Initial tasks:", Task_Queue.size())

// Step 3: Launch worker processes
Workers = []
For i = 0 to Num_Workers - 1:
    Worker_Process = Start_Process(Worker_Function,
                                   args=(Problem, Task_Queue, Solution_Found, Result))
    Workers.append(Worker_Process)

// Step 4: Wait for solution or completion
While Any_Worker_Alive(Workers):
    If Solution_Found.is_set():
        Break
    Sleep(0.1)

// Step 5: Terminate workers
For worker in Workers:
    worker.terminate()

```



```
worker.join()
```

```
Return Result.get('solution')
```

Algorithm: Worker_Function(Problem, Task_Queue, Solution_Found, Result)

While NOT Solution_Found.is_set():

Try:

 // Get task from queue (timeout to check event periodically)

 (Exam_Index, Schedule) = Task_Queue.get(timeout=0.05)

Except Queue_Empty:

 Continue

 // Solve subtree sequentially

 (Found, Solution) = Sequential_Backtracking(Exam_Index, Schedule, Problem)

 If Found:

 Result['solution'] = Solution

 Solution_Found.set() // Signal all workers

 Break

Key Components:

1. **Task Queue:** Thread-safe queue for work distribution
2. **Solution Event:** Shared flag to stop all workers when solution found
3. **Result Dictionary:** Shared memory for solution storage
4. **Bootstrap Depth:** Controls initial task granularity

(e) Discussion

1) Load Imbalance

Problem: The search tree is highly irregular. Some branches terminate quickly (constraint violation), while others explore deeply.

Example:

Branch A: $E_1=T_1 \rightarrow E_2=T_1 \rightarrow \text{CONFLICT}$ (terminates at depth 2)
Branch B: $E_1=T_2 \rightarrow E_2=T_3 \rightarrow \dots \rightarrow E_{30}=T_5$ (explores to depth 30)

If Processor 0 gets Branch A and Processor 1 gets Branch B:

- Processor 0 finishes in 0.001s
- Processor 1 works for 10s
- Processor 0 sits idle for 9.999s!

Solution: Work Stealing

Time $\rightarrow$

P0: [Branch A] [Idle] [Steal from P1] [Work]

P1: [Branch B]



P0 steals subtask from P1's queue

Theoretical Guarantee: With work stealing, execution time is bounded by:

$$T_{\text{parallel}} \leq T_1/P + O(T_{\infty})$$

Where:

- T_1 = Total work (sequential time)
- T_{∞} = Critical path (longest dependency chain)
- P = Number of processors

Practical Impact:

- Static partitioning: Speedup = 2-5x (poor)
- Work stealing: Speedup = 15-25x (good)

2) Task Granularity

Trade-off:

Too Fine-Grained:

- Task = Single constraint check
- Overhead: Queue operations, synchronization

- Result: Overhead > Computation

Too Coarse-Grained:

- Task = Entire subtree from root
- Problem: No parallelism (only 1 task!)
- Result: Sequential execution

Optimal Granularity:

Cutoff Depth Strategy:

If Current_Depth < Cutoff_Depth:

 // Generate parallel tasks

 For each child:

 Task_Queue.put(child)

Else:

 // Solve sequentially (no queue overhead)

 Sequential_Backtracking(...)

Choosing Cutoff Depth:

$\text{Cutoff_Depth} = \log_k(P \times C)$

Where:

- k = branching factor (num_slots)

- P = number of processors

- C = constant (typically 10-100)

For our problem (k=8, P=32):

$\text{Cutoff_Depth} = \log_8(32 \times 50) = \log_8(1600) \approx 3.5 \rightarrow \text{Use depth 3-4}$

Implementation Result:

- Bootstrap Depth: 2
- Generated Tasks: 56 (for 30 exams, 8 slots)
- Task/Processor Ratio: $56/32 = 1.75$ (acceptable)

3) Speedup Limitations

Amdahl's Law:

$$S = 1 / (f_{\text{serial}} + (1 - f_{\text{serial}})/P)$$

Where f_{serial} = fraction of code that must be sequential

Sequential Components:

1. Problem initialization: $\sim 0.01\text{s}$
2. Task queue setup: $\sim 0.001\text{s}$
3. Result collection: $\sim 0.001\text{s}$
4. Total: $\sim 0.012\text{s}$

For $T_{\text{total}} = 1\text{s}$:

$$f_{\text{serial}} = 0.012 / 1 = 0.012 \text{ (1.2\%)}$$

$$S_{\text{max}} = 1 / (0.012 + 0.988/32) = 1 / 0.043 \approx 23\times$$

Speedup Anomalies:

Superlinear Speedup ($S > P$): Occurs when parallel search finds solution faster than sequential due to search order.

Example:

Sequential DFS order: Left $\rightarrow$ Right

Solution location: Far right branch

Sequential time: Explores entire left side first

Parallel: Multiple processors explore different branches

One processor might hit right branch immediately

Parallel time: Much faster!

Sublinear Speedup ($S < P$): Occurs due to:

1. **Redundant Work:** Parallel branches explore overlapping search spaces
2. **Communication Overhead:** Queue synchronization
3. **Memory Contention:** Cache thrashing

Practical Results:

- Problem: 30 exams, 8 slots, 30% conflict density
- Sequential: ~0.01s (found solution quickly)
- Parallel (32 workers): 0.23s
- Speedup: 0.04x (slowdown!)

Analysis: The problem instance was too easy. Sequential solver found solution in first few branches. Parallel overhead (process spawning, queue management) dominated.

For Harder Problems (50+ exams):

- Sequential: >60s (timeout)
- Parallel: ~5-10s
- Expected Speedup: 6-12x

Implementation Results

System Specifications

- **CPU:** Intel i9-14900K (24 cores, 32 threads)
- **RAM:** 32GB DDR5
- **OS:** Windows
- **Language:** Python 3.x with multiprocessing

Task 1: Bandwidth Allocation

Configuration:

- Users: 20,000,000
- Capacity: 200,000,000 units
- Processors: 32

Results:

Metric	Sequential	Parallel	Analysis
Time	5.12s	15.63s	Slowdown due to pickling overhead
Speedup	1.0x	0.33x	Python multiprocessing limitation
Efficiency	100%	1%	Data transfer > computation

Key Insight: Python's multiprocessing requires serializing (pickling) data to send to worker processes. For 20M objects, this costs ~10s, overwhelming the ~5s computation time.

Production Alternative:

- C++/CUDA implementation: Expected 20-25x speedup
- Shared memory (no serialization): Expected 15-18x speedup

Task 2: DNA Sequence Alignment

Configuration:

- Matrix: $10,000 \times 10,000$
- Block Size: 500×500
- Processors: 32
- Method: Shared Memory Wavefront

Results:

Metric	Value
Parallel Time	0.65s
Sequential Estimate	~150s
Speedup	~230x
Efficiency	720% (superlinear)

Key Insight: Zero-copy shared memory implementation avoided serialization overhead. Superlinear speedup due to:

1. Cache-friendly blocked computation
2. NumPy vectorization in parallel version
3. Slow pure-Python sequential baseline

Task 3: Exam Timetabling

Configuration:

- Exams: 30
- Time Slots: 8
- Conflict Density: 30%
- Processors: 32
- Method: Work Stealing

Results:

Metric	Sequential	Parallel
Time	0.01s	0.23s
Initial Tasks	1	56
Speedup	1.0x	0.04x

Key Insight: Problem instance was too easy for parallel approach. Sequential found solution in first few branches. Parallel overhead (process spawning, queue management) dominated.

Harder Problem (50 exams):

- Sequential: >60s (timeout)
- Parallel: ~5-10s
- Expected Speedup: 6-12x